

# Advanced Networks

(EENGM4211)

Rasheed Hussain

[rasheed.hussain@bristol.ac.uk](mailto:rasheed.hussain@bristol.ac.uk)

Best contact : Microsoft Team

---

bristol.ac.uk

# Part 6: Quality of Service

# Unit Outline

- Introduction
- Internet Routing and Switching
- IP Multicast
- Networking for Realtime Applications
- Routing in Wireless Networks
- **Quality of Service**

# Outline & Topics

- Motivation
- Some practical principles
- How do we apply the principles to make QoS work?
- Case study : IP Quality of Service
  - What? why?
  - Protocols

# Motivation

---

## QoS Definition

From the end **user** point of view

QoS is a quantification of service/application-relevant measures of network effectiveness against acceptable levels of measures such as:

- Delay
- Jitter
- Packet Loss Rate
- Response time
- Throughput
- Service Availability

QoS from the **service providers** point of view:

Provide services to specific traffic classes such that QoS can be provided to end-users on a guaranteed/differential basis

→ Ability to provide better service to selected traffic

# Different categories of traffic

- Low latency (e.g. streaming)
- Guaranteed delivery (e.g. e-commerce)
- Best effort (e.g. email)

**Table 1** Internet Traffic Characteristics

<i>IntServ Category [1]</i>	<i>IntServ Subcategory</i>	<i>3GPP Category [2]</i>	<i>Examples</i>	<i>Constant Rate Necessary</i>	<i>Low Delay Necessary</i>	<i>Low Jitter</i>	<i>Low Delay Preferred</i>
Tolerant Real-Time		Streaming	Audio/video streaming	✓	✓		
Intolerant Real-Time		Conversational	IP telephony, teleconferencing	✓	✓	✓	
Elastic (Non-Real-Time)	Interactive Burst	Interactive	Telnet, HTTP				✓
	Interactive Bulk Transfer		FTP				✓
	Asynchronous Bulk Transfer	Background	SMTP				
...							

# Motivation

## Why should we care for QoS?

- Because...
  - Customers have expectations,... and **receive bills**
  - Operators must fulfil those expectations to justify the bills
  - It is an interesting theoretical and practical problem
- Historically there was one service – voice, wired up per call
  - Easy to manage QoS – certain number of calls per hour, lasting so many minutes gives number of circuits needed
  - Other things easier with digital (codecs, tones, not pulses, etc.)
- Then there was fax – pretty much the same as voice
- Then there was packet switching – the industry tried to pretend this was just like voice too, but it isn't
- Then the Internet got in the way...

# Motivation

---

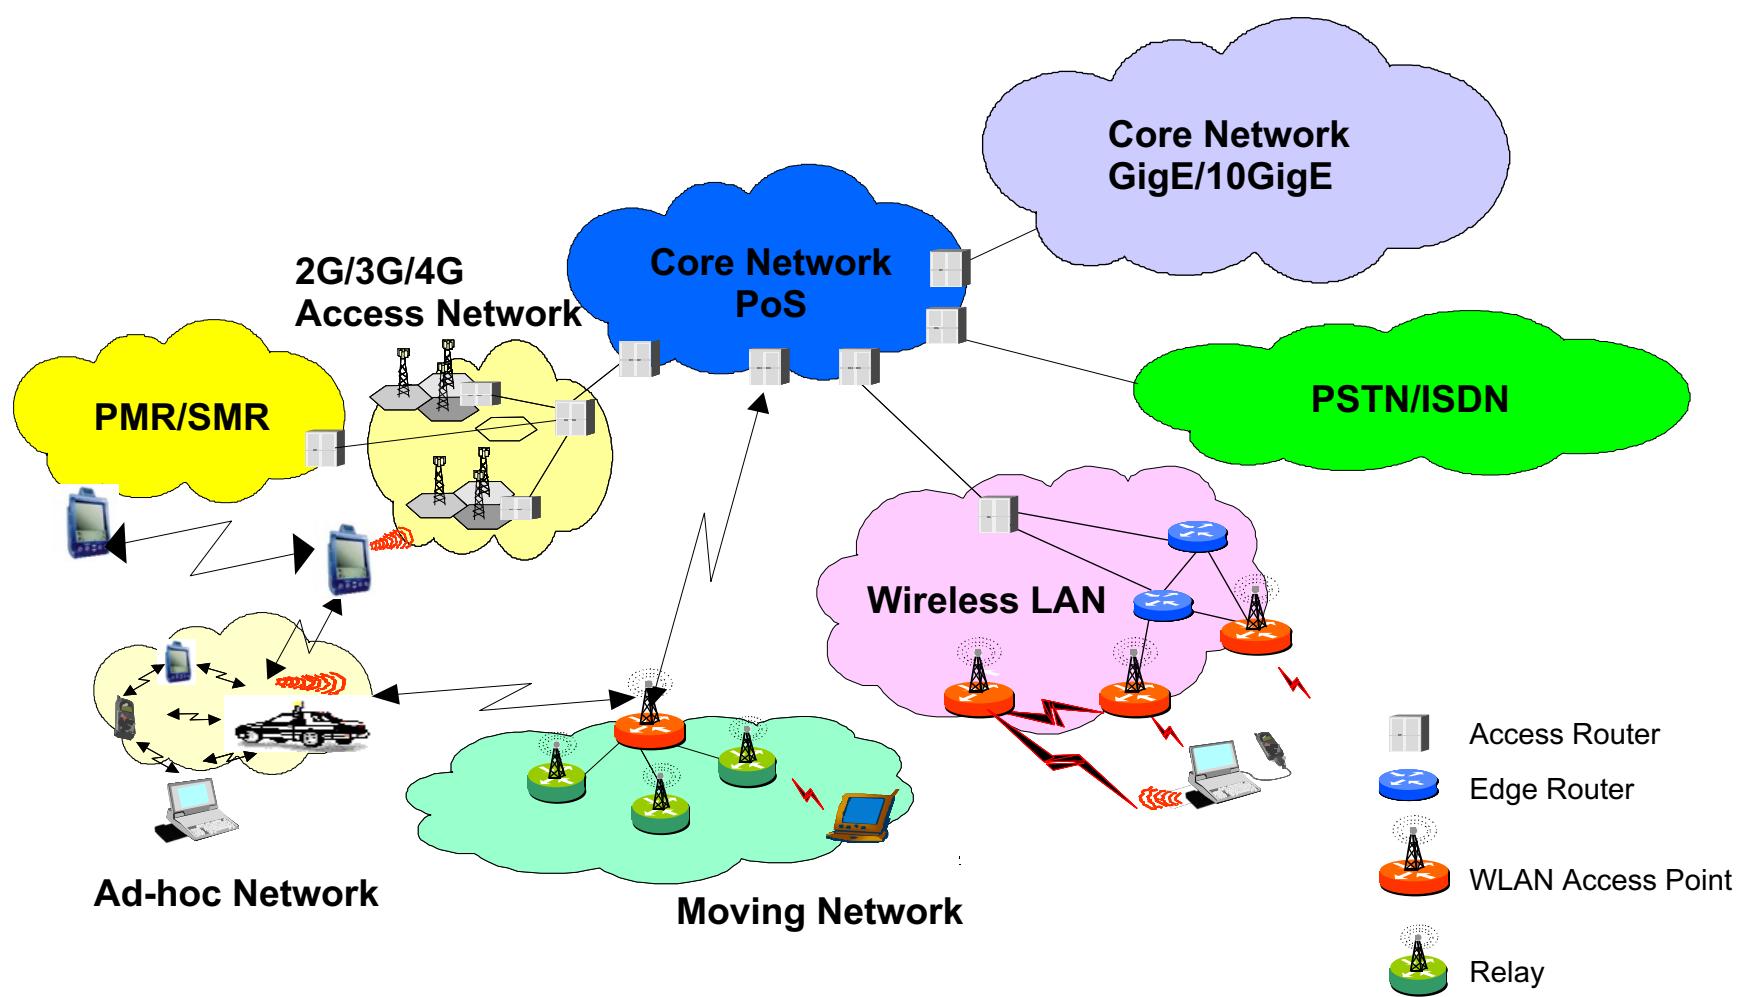
## Benefits of using QoS

- Traffic gets to destinations quickly and reliably to meet application requirements
- Controlled network latency and jitter and improved loss characteristics
- Effectively manage existing bandwidth and allocate resources in the network (bandwidth, buffer, priority etc)
- Control and predictability beyond “**best effort**” concept

# Section 1

Defining the problem  
Managing flows and resources

# Let's look at the Internet



# QoS: the problem

---

- QoS is only a problem when the communications system is **oversubscribed**
  - Guaranteed QoS:
    - If you know the whole state of the system over time
      - This isn't possible
    - Events and flows are predictable
      - No errors, no congestion
      - But... they aren't and we can't
    - Applications and users behave
  - Question: How can we improve the assurance that an application will receive a requested throughput, delay, error rate?
- 

# QoS: some questions

---

Why can't it be done in the best effort Internet?

Because:

- Routes change all the time.
- Traffic cannot be predicted or planned.
- Routers/switches and links are congestion points that cannot be controlled or managed.

How can this be fixed?

- end-to-end transport protocol guarantee delivery but cannot control delay or throughput
- Can asynchronous packet-switching technology provide guarantees of service? How?
  - e.g. OSPF takes into account the link load.

Why is this not sufficient?

- Response time; risk of instability in the network routing, etc.

# Internet Principles

- Global management framework for naming and addressing
- Common, limited, set of internetworking functions
- Local optimisations within a global framework
  - Backbone vs. border vs. local area – to make routing work
  - Communications layers – to separate concerns
- End-end principle
  - Functions placed at low levels of a system may be redundant when compared with the cost of providing them
  - Such functions cannot be completely and correctly implemented except with the knowledge and help of the application executing at the end points of the communication system.
- Fate-sharing principle
  - An end-to-end protocol design should not rely on the **maintenance of state** (i.e. information about the state of the end-to-end communication) inside the network.
  - Such state should be **maintained only in the end-systems**, in such a way that the state can only be destroyed when the end-system itself breaks.

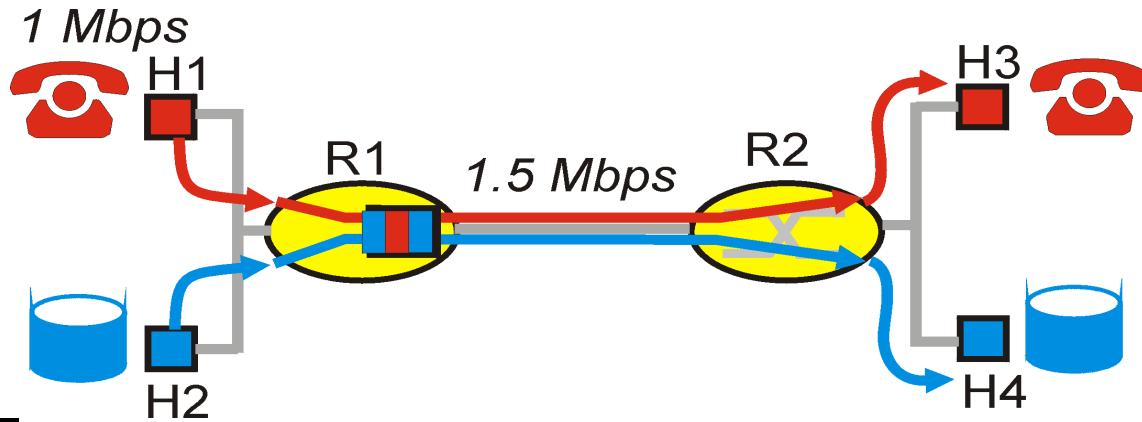
# Definitions

- A **flow**: is a sequence of packets sent between a source/destination host pair that follow the same network path.
- Allocate bandwidth: (on output links) the proportion of the link capacity used by a **flow**
- Manage queues of packets according to allocation of resources
  - Throughput, loss, delay
- Signals between routers
  - Flow specification
  - Traffic specification
  - Service requests

# QoS Principle 1 / 4: Use a marking and servicing protocol

Example: 1Mbps IP phone and an FTP connection share 1.5 Mbps link.

- FTP sender does not know what maximum rate is tolerable → bursts of FTP data traffic can congest router and cause audio loss
- Want to give priority to audio over FTP → how can the router tell which application a packet belongs to?

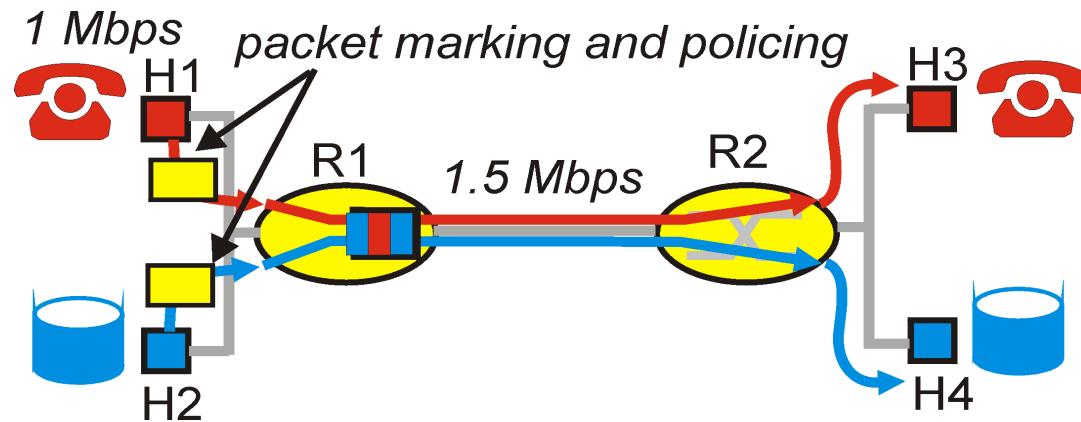


## Principle 1

*Packet Marking* needed for router to distinguish between different classes; and new router policy to treat packets accordingly

# QoS Principle 2 / 4: Policing: Make flows behave

- What if applications misbehave (audio sends higher than declared rate)
  - policing: force source adherence to bandwidth allocations
- Marking and policing at network edge



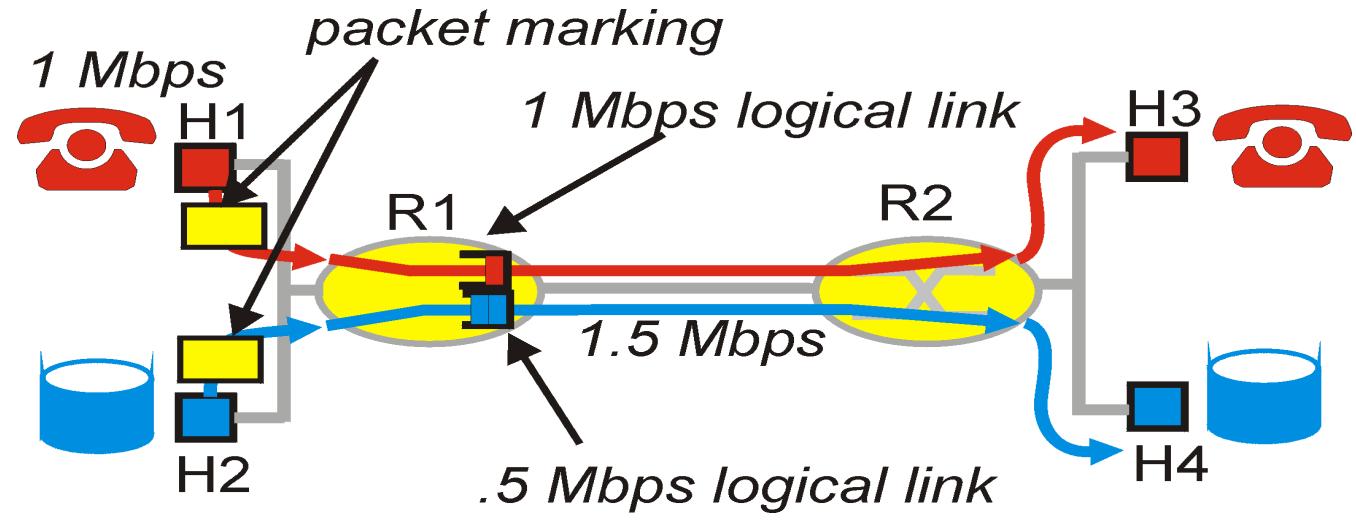
## Principle 2

Provide protection (*isolation*) for one flow from others

# QoS Principle 3 / 4: Maximise efficiency

Allocating *fixed* (non-sharable) bandwidth to flow:

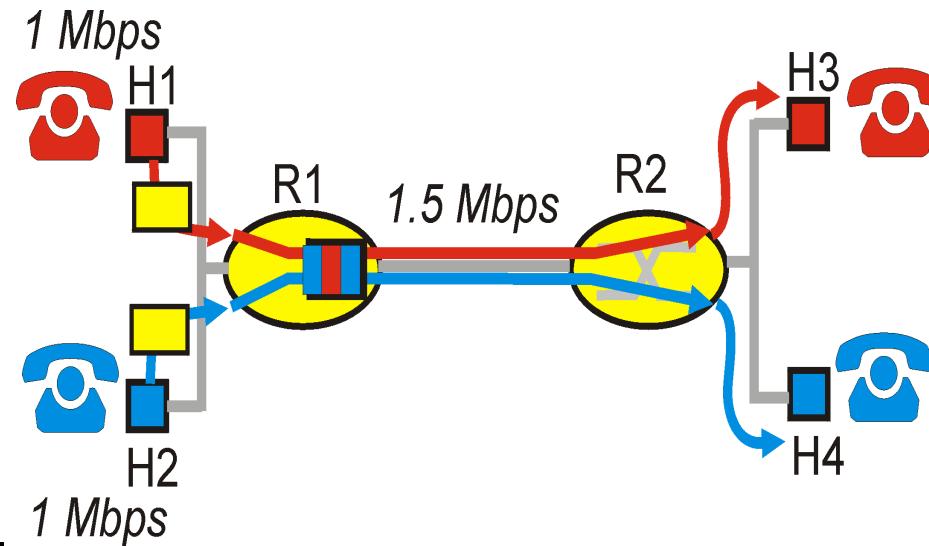
- *Inefficient* use of bandwidth if flows don't use their allocations



While providing isolation, it is desirable to **use resources as efficiently as possible**

# QoS Principle 4 / 4: Don't over-allocate

Fact: cannot support traffic demands beyond link capacity!



## Principle 4

*Call Admission:* A flow declares its needs. The network may block the call (e.g. busy signal) if it cannot meet the requirements.

# Basic functional blocks of QoS

## ■ Pieces of the Puzzle

### ■ Call Admission/ Resource Allocation

- SLA (Service Level Agreement)
- Advertisement of required bandwidth by user to the network
- Provisioning of bandwidth by network for user

### ■ Flow Classification

- Ability of the network to identify incoming packets (flows) and assign them a pre-defined level of service

### ■ Policing and Shaping

- Ability of the network to monitor flows to ensure conformance to advertised traffic characteristics and provisioned resources

### ■ Congestion Avoidance Mechanisms

- Ability to monitor buffer utilization levels and regulate flow rates to alleviate congestion on network links

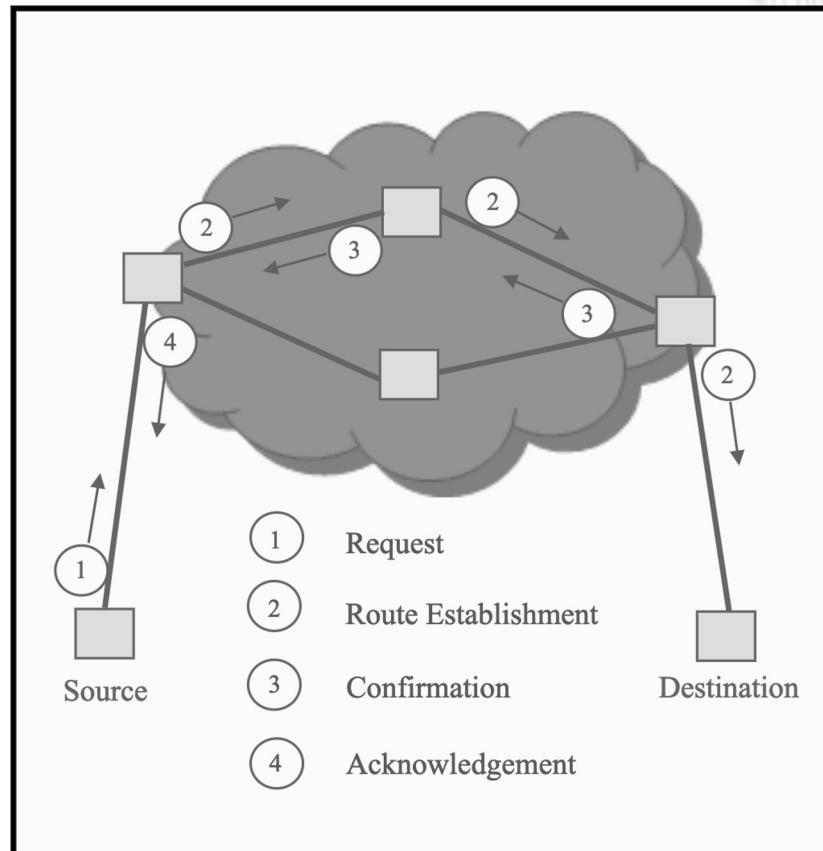
### ■ Scheduling Mechanisms

- Queueing mechanisms to provision differing levels of service

# Call admission

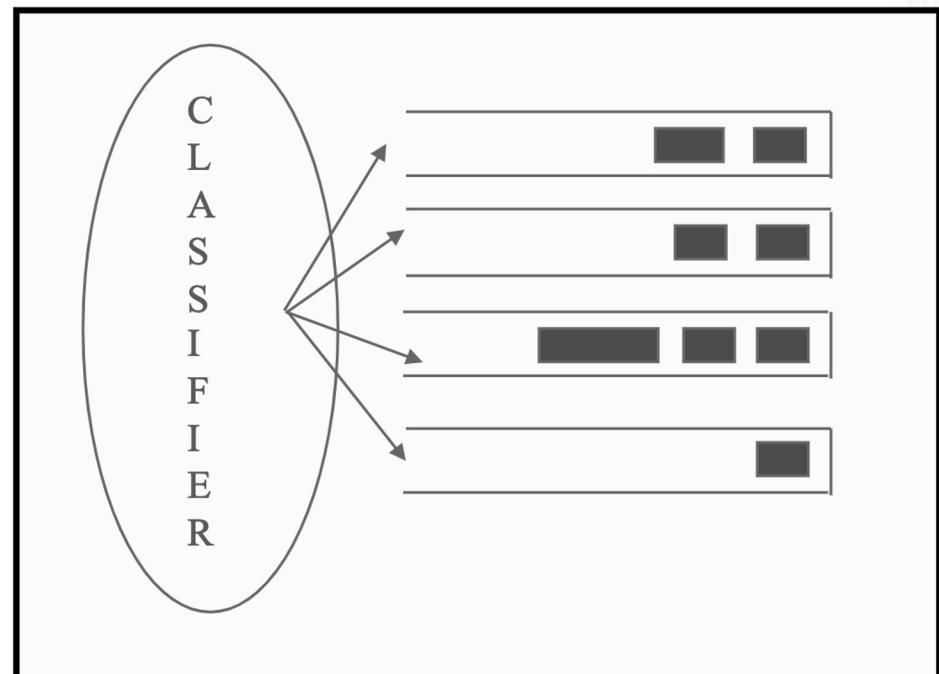
- Control Plane

- Signal the network on type of connection and QoS requirements
- Network is responsible for pro-active Bandwidth Management
  - Establishing the route of the connection
  - Reserving enough resources to meet QoS requirements
    - Stochastic reservation
    - Virtual pipes
  - Rejecting a call if it does not have enough resources to meet the call
- Examples:
  - ATM
  - IntServ (RSVP)
  - MPLS



# Flow Classification

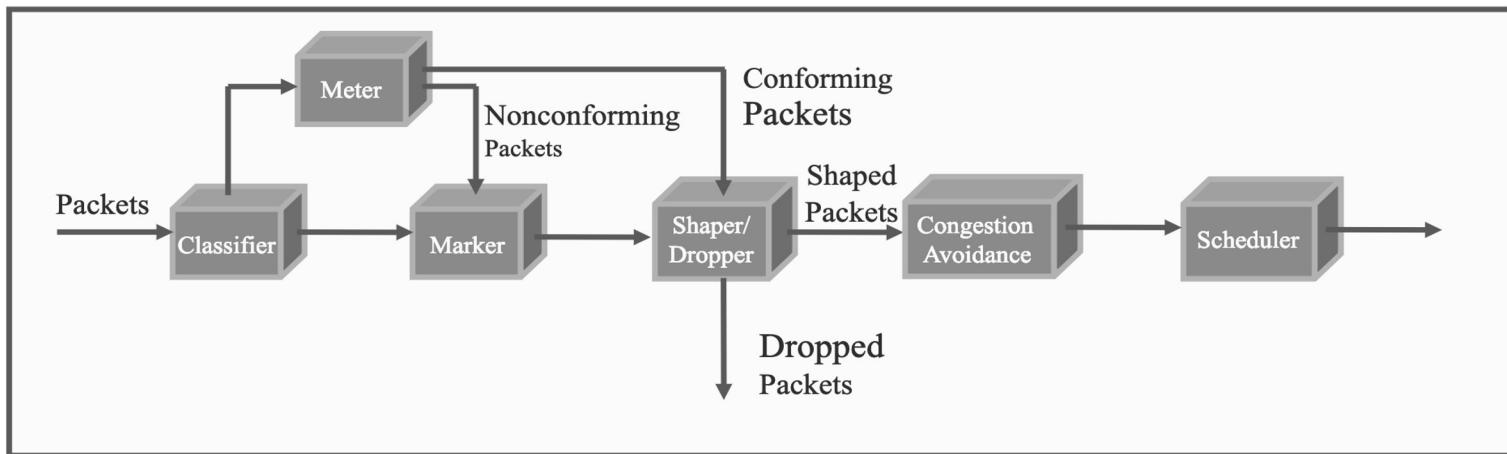
- *Need a method to identify packets/cells in order to provide differential treatment*
- Examples of Classification Criteria
  - VC number
  - MPLS Label
  - Type of Service
  - Protocol
  - Address
    - Source IP Address
    - Destination IP Address
  - Port Number
    - Source port
    - Destination port
  - Incoming interface



# Section 2

Ways of implementing QoS  
General purpose approaches

# Techniques for QoS



- **Classifier:** Selects packets based on portions of packet header
- **Marker:** Marks/Remarks the packet header based on traffic class
- **Meter:** Checks compliance to traffic profile and passes result to Marker and Shaper/Dropper
- **Shaper:** Allows for delaying of packets in buffer to enforce compliance with traffic profile
- **Dropper:** Drops traffic that does not conform with traffic profile
- **Congestion Avoidance:** Checks buffer levels and stochastically drops packets
- **Scheduler:** Allows for differential queueing and servicing of packets

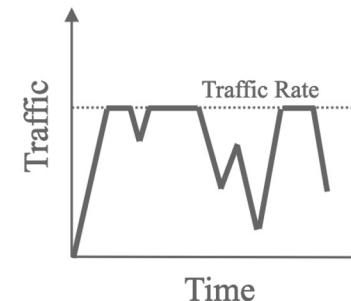
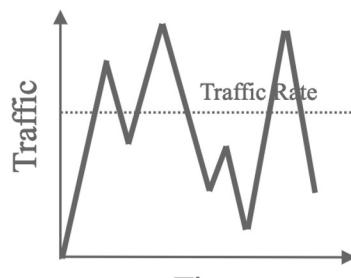
# Policing and Shaping

---

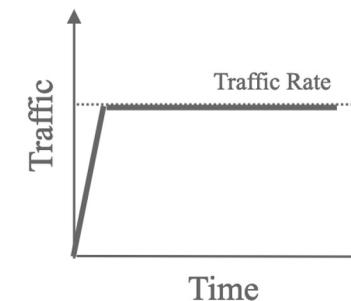
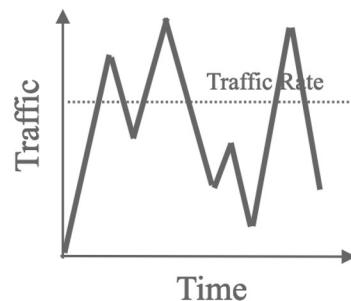
- *Regulate the rate at which a flow is allowed to inject packets into the network*
- Using an appropriate time scale for measurement is important for three metrics
  - Average rate
    - Limit the long term-average rate (packets per time interval) of the flow
  - Peak rate
    - Limit the maximum number of packets that can be allowed into the network over a short interval of time
  - Burst size
    - Limit the number of packets that can be allowed into the network over an extremely short interval of time
    - Limiting case (as time interval approaches zero) defines the number of packets that can be instantaneously sent into the network

# Policing and Shaping

- Effect of Policing



- Effect of Shaping



# Scheduling problem

---

- Assume we have  $K$  flows, each with its own queue.
- Problem:
  - (vers.1) In what order and for how long do I serve queue  $k_i$ ?
  - (vers.2) In what order and for how many bits I serve from queue  $k_i$ ?
  - (vers.3) In what order and when do I serve a packet from queue  $k_i$ ?
- Scheduling solutions:
  - FIFO (FCFS)
  - Simple priority queuing (also called strict priority)
  - Round-robin, Weighted Round-robin & Deficit Round-robin
  - Weighted Fair Queuing (WFQ)

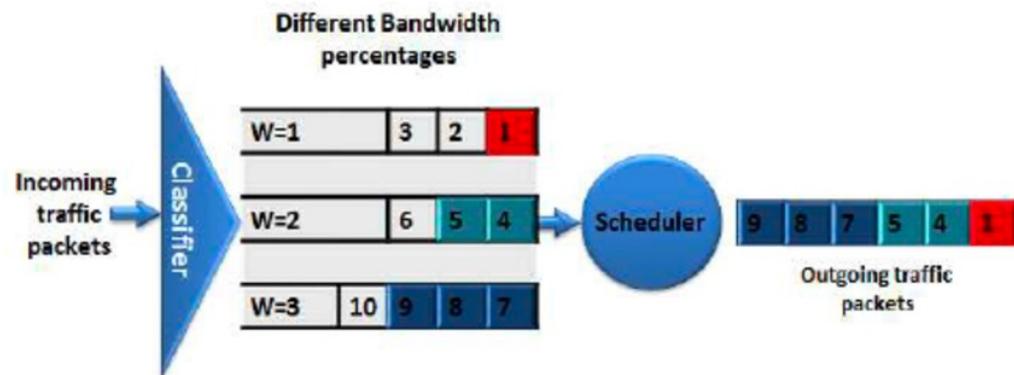
# Simple priority queuing

---

- $K$  queues:
  - $1 \leq k \leq K$
  - queue  $k + 1$  has greater priority than queue  $k$
  - higher priority queues serviced first
- ✓ Very simple to implement
- ✓ Low processing overhead
- Relative priority:
  - no deterministic performance bounds
- ✗ Fairness and protection:
  - not max-min fair: starvation of low priority queues

# Weighted round-robin (WRR)

- Simplest attempt at **generalized processor sharing (GPS)**
- Queues visited round-robin in proportion to weights assigned
- Different mean packet sizes:
  - weight divided by mean packet size for each queue
- Mean packets size unpredictable:
  - may cause unfairness



- Service is fair over long timescales:
  - must have more than one visit to each flow/queue
  - short-lived flows?
  - small weights?
  - large number of flows?

# Deficit round-robin (DRR)

- DRR does not need to know mean packet size
- Each queue has deficit counter (dc): initially zero
- Scheduler attempts to serve one quantum of data from a non-empty queue:
  - packet at head served if  $\text{size} \leq \text{quantum} + \text{dc}$   
 $\text{dc} \leftarrow \text{quantum} + \text{dc} - \text{size}$
  - else  $\text{dc} += \text{quantum}$
- Queues not served during round build up “credits”:
  - only non-empty queues
- Quantum normally set to max expected packet size:
  - ensures one packet per round, per non-empty queue
- Works best for:
  - small packet size
  - small number of flows

If weights are assigned to the queues, then the quantum size applied for each queue is multiplied by the assigned weight.

- Principle:
  - Model of continuous bit-by-bit transmission to produce finish-numbers for packets in queue
  - Serves packets round-robin
- Finish-number:
  - the time packet would have completed service under bit-by-bit scheduling
  - packets tagged with finish-number
  - smallest finish-number across queues served first
- Round-number:
  - execution of round by bit-by-bit round-robin server
  - finish-number calculated from round number
- If queue is empty:
  - finish-number is:  
number of bits in packet + round-number
- If queue non-empty:
  - finish-number is:  
highest current finish number for queue +  
number of bits in packet

# Weighted Fair Queuing (WFQ)

$$F(i, k, t) = \max \{F(i, k - 1, t), R(t)\} + P(i, k, t)$$

$F(i, k, t)$ : finish - number for packet  $k$   
on flow  $i$  arriving at time  $t$

$P(i, k, t)$ : size of packet  $k$  on flow  $i$   
arriving at time  $t$

$R(t)$ : round - number at time  $t$

$$F_\varphi(i, k, t) = \max \{F_\varphi(i, k - 1, t), R(t)\} + \frac{P(i, k, t)}{\varphi(i)}$$

$\varphi(i)$ : weight given to flow  $i$

- Rate of change of  $R(t)$  depends on number of active flows (and their weights)
- As  $R(t)$  changes, so packets will be served at different rates

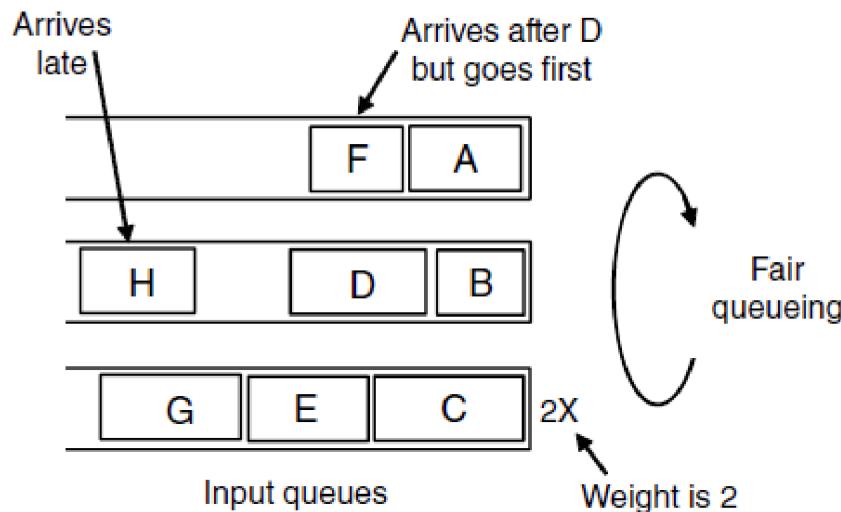
- Flow completes (empty queue):
  - one less flow in round, so
  - $R$  increases more quickly
  - so, more flows complete
  - $R$  increases more quickly
  - etc. ...
  - **Iterated deletion** problem
    - Transmit avalanche ....
- WFQ needs to evaluate  $R$  each time packet arrives or leaves:
  - processing overhead

# Weighted Fair Queuing (WFQ)

---

- WFQ is the only scheduler that also proposes packet dropping policy
- Buffer drop policy:
  - packet arrives at full queue
  - drop packets already in queue, in order of decreasing finish-number
- Can be used for:
  - best-effort queuing
  - providing guaranteed data rate and deterministic end-to-end delay
- WFQ used in “real world”
- Alternatives also available:
  - self-clocked fair-queuing (SCFQ)
  - worst-case fair weighted fair queuing (WF2Q)

# Weighted Fair Queuing

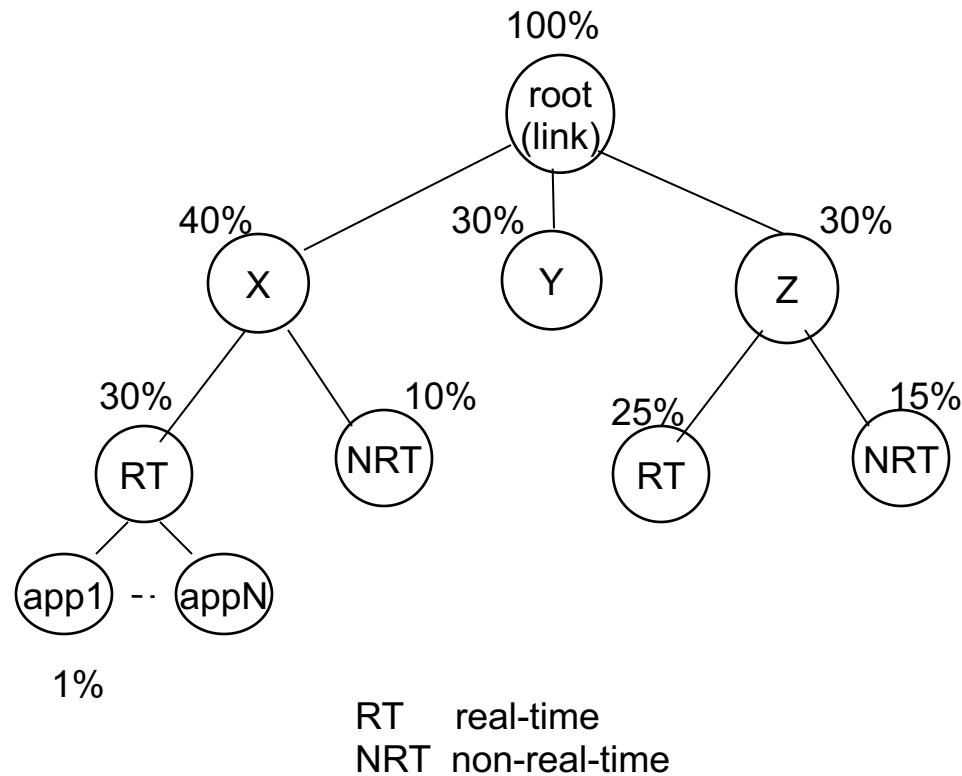


Packet	Arrival time	Length	Finish time	Output order
A	0	8	8	1
B	5	6	11	3
C	5	10	10	2
D	8	9	20	7
E	8	8	14	4
F	10	6	16	5
G	11	10	19	6
H	20	8	28	8

$$\text{Finish}(i) = \max(\text{Arrive}(i), \text{Finish}(i-1)) + \text{Length}/\text{Weight}$$

# Class-Based Queuing

- Hierarchical link sharing:
  - link capacity is shared
  - class-based allocation
  - policy-based class selection
- Class hierarchy:
  - assign capacity/priority to each node
  - node can “borrow” any spare capacity from parent
  - fine-grained flows possible
- Note: this is a queuing mechanism: requires use of a scheduler



# Summary

---

- Quality of Service
- Scheduling
- Queuing
- Policing
- Classification
- Scheduling methods